

PACT: Evidence-Bound Planning for Contract-Governed Tool Agents

Pu Yang

Independent Researcher
npuyangpu@gmail.com

Abstract

Tool-using language agents must interpret under-specified requests while respecting mutable capability contracts and avoiding unsupported claims of success. We present PACT, a CAR-bench Track 2 agent that restricts its Cerebras-hosted language model to semantic proposal. The model compiles the current policy, request, capabilities, and evidence into a bounded obligation graph; a deterministic kernel alone may release the next user interaction or tool call. Before release, the kernel strictly decodes the graph, validates its structure and arguments against live Draft 2020–12 schemas, and enforces capability freshness, one-shot confirmation scopes, and terminal evidence requirements. PACT recompiles when an event introduces a new semantic binding; otherwise a control-only event advances an immutable verified suffix. Every Act result reported as successful remains a completion obligation across later horizons. Each horizon uses one proposal, at most one verifier-guided repair, and, for action-bearing plans, one separate same-model audit. We state conditional guarantees for capability confinement, scoped confirmation, and evidence-closed completion; language interpretation remains outside the assurance boundary. Failure-oriented tests exercise rejection, contract drift, event mismatch, scope reuse, and carried-obligation violations; digest-pinned A2A validation establishes packaging and telemetry, not task quality.

1 Introduction

CAR-bench exposes brittle control through multi-turn state changes, missing capabilities, and ambiguous requests [Kirmayr *et al.*, 2026]. Its A2A evaluator sends policy, conversation, live tool declarations, and tool results [A2A Protocol Working Group, 2026]; the agent selects the next interaction, while the evaluator executes tools. Track 2 additionally audits sequential model-call depth and token telemetry.

The reliability problem is one of authority. A conventional reason–act loop lets one stochastic component interpret the request, construct arguments, accept confirmation, and declare completion. Valid JSON constrains syntax but does not show

that an operation remains live, that its arguments satisfy the current schema, or that a result justifies completion.

PACT separates semantic flexibility from execution authority. The model proposes a short obligation graph; deterministic code admits and advances it only after checking live contracts, dependencies, confirmation, and evidence. An Observe result or Ask reply is a semantic *data event* and starts a new horizon. An exact confirmation or matched Act success is a *control event* and advances the verified suffix without another model call. Every Act record reported as successful remains mandatory completion evidence across later horizons.

Contributions. PACT makes two methodological contributions. First, *verified-horizon execution* lets events that only discharge a checked obligation advance an immutable suffix, while events that may change an operation, argument, or scope invalidate the horizon and re-enter compilation. Second, *cross-horizon completion closure* keeps every reported-success Act record mandatory for any later completed outcome. The implementation additionally exposes a bounded serial inference graph (Figure 1) and a deterministic validation suite. The kernel contains no task IDs, benchmark utterances, operation allowlists, or domain workflows.

2 Related Work and Positioning

ReAct interleaves model reasoning and acting [Yao *et al.*, 2023]. LLMCompiler optimizes parallel function dispatch [Kim *et al.*, 2024]; PlanCompiler validates typed pipelines over a fixed primitive registry [Harikumar, 2026]; and AgentSpec externalizes runtime rules [Wang *et al.*, 2026]. As scoped here, none couples a revocable, per-emission capability contract to terminal evidence carried across planning horizons. PACT instead serializes visible obligations, binds A2A transitions to reported outcomes, and revalidates the live contract before release.

The producer/checker split is inspired by proof-carrying code [Necula, 1997] and runtime verification [Leucker and Schallhart, 2009], but PlanIR is a locally checked execution artifact, not a proof of language meaning.

3 PACT

3.1 Trust Model and State

PACT trusts its sole-emission adapter, PlanIR/schema validators, digest checks, and evaluator policy/capability metadata.




Figure 1: One Track 2 decision horizon, corresponding to a baseline LLM step. The maximum completed-model path is proposal → repair → Act audit: three serial calls. The audit is used only for an accepted Act plan, and its replacement is fully reverified. A rate-limited transport attempt yields no completion or tokens; exhaustion fails safely. Ask/Observe data start a new horizon, whereas exact confirmation and matched Act success advance the immutable suffix.

Result values remain untrusted semantic data, but the claims assume that the evaluator returns the result of the sole pending call and truthfully reports its explicit SUCCESS/FAILURE status. PACT validates the envelope and local binding; it does not independently attest the external effect. A2A supplies an evaluator call ID but echoes no agent-selected nonce.

For each active goal, the adapter stores the live contract T_t , an append-only evidence ledger E_t , the accepted graph G_t , at most one waiting runtime obligation p_t , reported-success Act-key obligations O_t , consumed confirmation scopes C_t , and accumulated inference metrics M_t . Values cross the ledger through a JSON round trip, preventing caller mutation of nested history. Event fingerprints make identical replay idempotent; conflicting reuse and out-of-order input cannot satisfy a waiting node. Per-context locks serialize turns, and cancellation deletes goal state without replacing its lock.

3.2 Compile, Verify, and Audit

The semantic compiler receives the trusted policy in a separate system partition and a bounded view of the goal, recent events, T_t , E_t , and O_t . Its provider-facing language is deliberately smaller than PlanIR. Because Cerebras strict response schemas forbid unconstrained object properties [Cerebras Systems, 2026], operation arguments travel as one `arguments_json` string. The local decoder rejects duplicate keys and non-finite numbers, decodes exactly one JSON object, and validates it against the live Draft 2020–12 operation schema [JSON Schema, 2022]. Provider constrained decoding is therefore an early filter, not the authority; empirical structured-generation work motivates retaining this local check [Geng *et al.*, 2025].

PlanIR is a bounded DAG over

{OBSERVE, ASK, CONFIRM, ACT, RESPOND}.

OBSERVE and ACT identify a live operation, constant arguments, and a fresh success-evidence key. ASK and CONFIRM wait for a user event; each CONFIRM explicitly names the

descendant Act nodes it authorizes. The sole terminal RESPOND has one of five outcomes: completed, answered, refused, declined, or fail-safe. The verifier checks unique identifiers, existing dependencies, acyclicity, terminal reachability, at most 20 nodes, dependency depth at most 12, evidence freshness, operation membership, and the complete recursive argument schema.

One rejected proposal may be repaired from value-redacted typed issues. A second rejection fails safely. After a plan containing an Act node passes locally, a separate same-model prompt audits the full policy, goal, evidence, capabilities, and argument provenance. The audit returns a complete replacement candidate, which crosses the entire decoder and verifier again; failure never falls back to the pre-audit plan. This pass is a stochastic cross-check, not a semantic proof. The submitted default is `gpt-oss-120b`, medium reasoning, and 8,192 maximum completion tokens. A provider `length_finish` is a typed truncation and does not trigger a futile identical semantic repair.

3.3 Execute, Replan, and Close Evidence

The runtime deterministically chooses the first topologically ready node, emits one runtime obligation, and waits. Before an external Observe or Act call, it rebuilds the complete capability snapshot, requires an unchanged digest, rechecks operation membership and any known effect label, and validates a deep-copied argument object again. A revoked or changed contract therefore cannot inherit admission from an older snapshot.

Verified-horizon principle. An event is *control-only* when it supplies no value that can parameterize an unexecuted node and merely discharges the unique waiting obligation. An exact positive confirmation and a reported-success Act result whose continuation depends only on status are control-only: PlanIR has constant arguments and no expression language, so neither event can rewrite the accepted suffix. If G_t was verified under digest d , the event matches p_t , and the live digest remains

d , processing changes only node status and the prescribed ledger, O_t , or C_t record. Names, arguments, dependencies, and remaining order are unchanged, so the next ready node may be released without semantic recompilation, subject to the emission guard.

Ask replies, Observe results, and qualified confirmations instead introduce a value or scope with no verified binding in the old graph. PACT discards that horizon, preserves E_t and O_t , and recompiles. This data/control cut is a validity boundary, not a latency heuristic. It preserves checked structure and evidence obligations, not semantic correctness. Each reported-success Act key enters O_t , so later replanning cannot silently forget it.

The adapter accepts exactly one external result, requires the pending operation name and a nonempty evaluator call ID, then binds the payload to the sole pending context, plan, node, argument digest, and capability digest. The runtime validates this internally bound event. Correlation is therefore unambiguous under the single-pending assumption but is not independently authenticated by echoed metadata. Only an explicit reported SUCCESS creates successful external evidence. Local lowering replaces model-proposed completion citations with the current Act keys plus O_t and checks their producers and status. An answered outcome contains no Act node; semantic effect coverage requires trusted effect metadata.

Confirmation is required only when trusted metadata sets `x-pact-requires-confirmation` or the trusted description begins with the anchored marker `REQUIRES_CONFIRMATION`. A qualifying authorization is either a same-plan ancestor or an unused ledger scope

$$(D(T_t), o, \text{SHA256}(x)), \quad (1)$$

where $D(T_t)$ is the digest of the entire capability snapshot, o the operation, and x its canonical arguments. A matching successful Act consumes only its individual scope. Using the full snapshot is conservative: even an unrelated capability change invalidates prior authorization.

4 Enforced Properties and Compute

Table 1 separates locally enforced properties from their assumptions. Three invariants carry the main argument. First, if PACT emits tool call (o, x) , then $o \in T_t$, $x \models \text{schema}_{T_t}(o)$, and the current snapshot digest equals the admission digest. This follows because admission and the sole emission path perform the same membership and schema checks.

Second, let $\mathcal{K}_A(E)$ be the evidence keys in E whose source is external, producer is ACT, and explicit status is `success`.

$$\text{completed} \implies \text{ActKeys}(G_t) \cup O_t \subseteq \mathcal{K}_A(E_t). \quad (2)$$

Unmatched, failed, duplicated, or conflicting events cannot satisfy a waiting node. Current keys originate at Act producers; carried keys must resolve to reported-success external Act records. Induction over horizons establishes obligation monotonicity: a reported-success Act adds one key to O_t , and no nonterminal replan removes it.

Third, confirmation requires a same-plan ancestor or unused scope from Equation 1; success consumes that scope and C_t prevents reuse.

Property	Enforcement and assurance boundary
Call confinement	Live membership and full Draft 2020–12 argument validation at admission and emission under one digest; assumes the evaluator’s declaration accurately describes the operation.
Transition integrity	One waiting obligation; operation and nonempty evaluator-ID checks followed by internal binding/validation. A2A provides no echoed agent nonce.
Scoped confirmation	Same-plan ancestry or unused $(D(T_t), o, \text{SHA256}(x))$ scope; individual consumption after Act success. Requires trusted metadata and confirmation parsing that reflects intended consent.
Completion closure	Reported-success Act evidence plus all carried keys. Real effect coverage requires truthful status and trusted effect metadata.
Compute audit	Depth at most three and fan-out one; transport attempts, tokens, and normally-returned-call quota wait are separated.

Table 1: The assurance boundary: what code enforces and what remains an environmental or semantic assumption.

Track 2 compute audit. Horizon j replaces one baseline-agent LLM call. With $r_j, a_j \in \{0, 1\}$ indicating repair and Act audit, its completed depth is

$$L_j = 1 + r_j + a_j \leq 3 \leq 5, \quad (3)$$

with fan-out one: no-Act branches use one or two completions and Act branches two or three. Each invocation permits one initial HTTP request plus at most three rate-limit retries, hence at most 12 attempts and 1,800 seconds of scheduled rate-limit wait per horizon. A rate-limited attempt returns no model completion or tokens; exhaustion raises a typed error.

Cerebras reports reasoning inside provider completion tokens. PACT therefore exports `prompt.tokens` as input tokens, `thinking.tokens` as reasoning tokens, and `completion.tokens` as provider completion minus reasoning. These fields are mutually exclusive and accumulate across rejected proposals, repair, audit, and tool turns until the next text response. For N tasks and their reported-response sets R_k , let p_r, h_r, c_r denote input, reasoning, and non-reasoning completion tokens. Then $U_k = \sum_{r \in R_k} (p_r + h_r + c_r)$, and Track 2 requires $N^{-1} \sum_k U_k \leq 500,000$ [CAR-bench Challenge, 2026]. PACT reports the counters but, for an open conversation, does not prove this inequality or global termination. `avg_llm_call_time.ms` averages successful SDK-request duration; `quota_wait_time.ms` sums quota-eligible wait only for normally returned invocations. Wait accumulated by a retry-exhausted invocation is absent from final telemetry. Neither metric is end-to-end A2A latency.

5 Release Validation

We validate mechanisms and the submitted artifact, not hidden outcomes. Table 2 summarizes failure-oriented coverage.

At report freeze on 23 July 2026, `uv run pytest -q` passed 356 tests and three subtests. Ruff passed the full tree, and mpy passed the 18-file release closure.

Mechanism	Adversarial coverage → release behavior
Plan admission	Duplicate IDs/keys, cycles, excessive depth, unreachable terminals, malformed wire JSON, exhausted repair, and rejected audit → no admission or tool call.
Live contracts	Nested schemas, combinators, additional properties, malformed declarations, effect-label mismatch, and revocation → local validation, replan, or emission-time rejection.
Event provenance	Failure, replay, conflicting event ID, wrong operation/digest, out-of-order result, and caller mutation → no success or state advance.
Consent and closure	Stale/reused scope, changed arguments, later failure/replan → per-scope consumption and prior Act keys remain mandatory.
Compute and protocol	Truncation, bounded 429s, cancellation, concurrency, and telemetry → bounded semantic/retry paths, serialized context, and typed exhaustion.

Table 2: Mechanism-validation matrix for the final configuration. These tests establish executable guard behavior, not CAR-bench reward or Pass³.

5.1 Failure Semantics and Validation Oracle

PACT treats failure as an absence of execution authority, rather than merely as negative text. Admission or audit rejection and provider-budget exhaustion create no runtime and release no external call. Contract invalidation discards G_t and recompiles with E_t and O_t unchanged; any later outcome passes the ordinary completion guard. A mismatched, replayed, conflicting, or out-of-order event cannot satisfy p_t . Negative confirmation and explicit external failure terminate the suffix without a new Act. No failure path falls back to an unaudited candidate or invents compensation. Set $\Delta_t^A = \mathcal{K}_A(E_{t+1}) \setminus \mathcal{K}_A(E_t)$; let $S(e_t) = \{\kappa\}$ for a valid, matched Act-SUCCESS event carrying key κ , and $S(e_t) = \emptyset$ otherwise. Let y_t be the emitted outcome. For typed failure f_t , the oracle is

$$\Delta_t^A \subseteq S(e_t), \quad f_t \implies y_t \neq \text{completed}. \quad (4)$$

This permits a valid Act success before a later closure failure; failure neither fabricates nor erases one. Scopes enter C_t only after matching success. Partial results are disclosed; no rollback is claimed because a contract may lack an inverse.

Oracle-driven protocol. Positive traces exercise valid advancement; paired negatives alter one invariant-bearing field—plan ID, operation, argument or capability digest, event type, or status—and exercise Equation 4. Compiler/verifier tests mutate reachability, evidence producers, confirmation ancestry, and carried obligations before runtime construction. Cross-horizon tests try to omit or relabel a carried reported-success Act key. Confirmation traces vary arguments, snapshots, and prior consumption; retry tests use a synthetic clock to reject a sleep before the cumulative budget is exceeded. Of 356 collected cases, 192 directly target the PACT kernel and 52 cover A2A response/scenario contracts; the remainder are not claimed as kernel evidence. These finite traces witness implemented guards, not exhaustiveness over all event sequences. They establish mechanism conformance; digest-pinned A2A

smoke establishes deployability. Neither estimates planning quality, audit efficacy, CAR-bench reward, or Pass³.

Necessity stress trace. Consider a plan admitted under digest d_0 , authorization for (o, x, d_0) , a reported-success Act, and a later replan that omits its key. Admission-only schema checking cannot stop release after the contract changes to d_1 ; a conversation-level confirmation flag can authorize changed arguments or snapshot; and a terminal check limited to the newest graph can forget the earlier Act. PACT assigns these failures to emission-digest equality, an exact one-shot scope, and O_t closure, respectively. PlanIR is therefore a short-lived, contract-bound execution certificate, while ledger obligations carry commitments across horizons.

Reproducible release identity. The public unprivileged `linux/amd64` image was built from source revision 3255 c9320bfa36e3fe0c4647ffc401b78e7259a6 and pinned as sha256:173fd630e1691b40af55e731e82023c1df9cfcae52d3b84420c1d307eccea6d1. An anonymous registry resolution succeeded. README Option C pulled this exact digest, received HTTP 200 from both Agent Cards, and all three public-format harness projections terminated at the A2A boundary with client exit code zero. It verified `turnmetrics` on nine final responses; no task success is inferred. The result artifact has SHA-256 55c43af4ec481a39c609554e2ef9fbe4315816d0340bb892f2f6689036115bbf. This is packaging, protocol, and telemetry evidence, not a task-quality estimate. We report no final-configuration public aggregate, hidden-set score, or ranking in the deadline version.

6 Limitations and Conclusion

PACT narrows execution authority but does not verify language meaning. Its same-model audit may reject or replace an Act candidate, but its error-reduction effect is unmeasured and it is not independent evidence. With `unknown` effect metadata, PACT cannot establish Observe-versus-Act intent. An Act result is control-only: the suffix may depend on status, not a new payload value; such workflows need an explicit observation boundary. The scope guarantee begins after deterministic confirmation parsing and does not prove intended consent. Single-pending A2A correlation is not cryptographic, the ledger is process-local, and safe failure may reduce coverage.

Within those boundaries, PACT gives a precise separation: the model proposes meaning, while deterministic code owns capability admission, evidence-bearing state transitions, scoped authorization, and the evidence conditions of a typed terminal outcome. Combining that boundary with verified-horizon execution yields a locally checkable, protocol-accounted Track 2 agent without benchmark-shaped workflows. The broader implication is that replanning assurance is not a property of a plan alone, but of a plan, its live contract, and a monotone evidence projection. This boundary transfers to other tool domains only when status, effect, and correlation assumptions are equally explicit.

References

- [A2A Protocol Working Group, 2026] A2A Protocol Working Group. Agent2Agent protocol specification, version 1.0.0. Linux Foundation, 2026. Accessed: 2026-07-17.
- [CAR-bench Challenge, 2026] CAR-bench Challenge. Competition tracks: Track 2 cerebras fast-reasoning. Challenge specification, 2026. Accessed: 2026-07-23.
- [Cerebras Systems, 2026] Cerebras Systems. Structured outputs. Cerebras Inference Documentation, 2026. Accessed: 2026-07-17.
- [Geng *et al.*, 2025] Saibo Geng, Hudson Cooper, Michał Moskal, Samuel Jenkins, Julian Berman, Nathan Ranchin, Robert West, Eric Horvitz, and Harsha Nori. Generating structured outputs from language models: Benchmark and studies, 2025.
- [Harikumar, 2026] Pranav Harikumar. PlanCompiler: A deterministic compilation architecture for structured multi-step LLM pipelines, 2026.
- [JSON Schema, 2022] JSON Schema. JSON Schema Core, draft 2020-12. Specification, 2022. Accessed: 2026-07-17.
- [Kim *et al.*, 2024] Sehoon Kim, Suhong Moon, Ryan Tabrizi, Nicholas Lee, Michael W. Mahoney, Kurt Keutzer, and Amir Gholami. An LLM compiler for parallel function calling. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 24370–24391. PMLR, 2024.
- [Kirmayr *et al.*, 2026] Johannes Kirmayr, Lukas Stappen, and Elisabeth André. CAR-bench: Evaluating the consistency and limit-awareness of LLM agents under real-world uncertainty. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 40599–40618, San Diego, California, 2026. Association for Computational Linguistics.
- [Leucker and Schallhart, 2009] Martin Leucker and Christian Schallhart. A brief account of runtime verification. *Journal of Logic and Algebraic Programming*, 78(5):293–303, 2009.
- [Necula, 1997] George C. Necula. Proof-carrying code. In *Proceedings of the 24th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 106–119, 1997.
- [Wang *et al.*, 2026] Haoyu Wang, Christopher M. Poskitt, and Jun Sun. AgentSpec: Customizable runtime enforcement for safe and reliable LLM agents. In *Proceedings of the 48th IEEE/ACM International Conference on Software Engineering*. ACM, 2026.
- [Yao *et al.*, 2023] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *Proceedings of the Eleventh International Conference on Learning Representations*, 2023.